

Raibert planner: the measured v_B does not beat the feedforward v_{cmd}

Both policies are 53-dimensional and were swept over the identical grid. Bars = mean of the five seed means, whiskers = std over seeds, hollow dots = the individual seed means. Δ is the v_B policy relative to the v_{cmd} policy. Both policies are evaluated at the same training iteration (checkpoint 14800), so training length is no longer confounded with the planner change.

